


Flow Based Models



James Iorio
Mentored by: Sebastian
Gutierrez-Hernandez



What's Covered

- Components of Flow Models
 - Couplings
 - Sigma
 - Paths
- Image Generation Example
- Challenges

Components of Flow Models: Coupling



How we define endpoint pairs for flow models


Components of Flow Models: Sigma

Overfitting vs. Underfitting

20000: loss 7.779 time 44.56




- Prior sample $z(S)$
- Flow
- $z(0)$




15000: loss 11.193 time 45.71




- Prior sample $z(S)$
- Flow
- $z(0)$




Exploration of Sigma



Exploration of Sigma





Components of Flow Models: Paths



Sample Image Generation

MLP vs. UNET Models



Hardest Part & What Made It Work?

```
# Define device: Use GPU if available, else CPU
use_cuda = torch.cuda.is_available() # Check if CUDA-capable GPU is available
device = torch.device("cuda" if use_cuda else "cpu")

# Training Parameters:
batch_size = 32 # Images per batch
num_epochs = 10 # Number of times to loop through full training dataset

# Load MNIST Dataset:
trainset = datasets.MNIST(
    "../data", # Folder where dataset is saved/downloaded
    train=True, # Loads training set
    download=True,
    transform=transforms.Compose([ # Converts PIL images to tensors and normalizes pixel values to [-1, 1]
        transforms.ToTensor(), # Convert PIL Image to PyTorch tensor
        transforms.Normalize((0.5), (0.5)) # Normalize grayscale channel to mean=0.5, std=0.5
    ]),
)

# Create a DataLoader to handle batching and shuffling
train_loader = torch.utils.data.DataLoader(
    trainset,
    batch_size=batch_size,
)

# 4. Training Loop:
num_epochs = 10
num_ode_steps = 5 # Number of time steps per trajectory

# Main Training Loop:
for epoch in range(num_epochs):
    for i, (x1, y) in enumerate(train_loader): # Loop over batches from DataLoader
        optimizer.zero_grad() # Reset gradients

        # Move to device and ensure 4D shape
        x1 = x1.to(device) # Target image batch
        y = y.to(device) # Conditional labels
        if x1.ndim == 3: # MNIST: (batch, H, W) -> (batch, 1, H, W)
            x1 = x1.unsqueeze(1)

        # Sample initial noise
        x0 = torch.randn_like(x1) # Same shape as x1

        # Sample discrete times for integration
        times = torch.linspace(0., 1., steps=num_ode_steps).to(device)

        # Initialize state with noise
        xt = x0

        # 3. CFM Class:
        sigma = 0.1 # Noise scale
        # Utility to convert tensors to images
        to_pil = ToPILImage()

        # Create Raw UNet Model:
        raw_unet = UNetModel(
            dim=(1,28,28), # (channels, height, width) of input images
            num_channels=32, # Base number of feature maps
            num_res_blocks=1, # Number of residual blocks per level
            num_classes=10, # Number of conditional classes (10 for MNIST)
            class_cond=True # Whether model is conditioned on labels
        )

        # Wrap UNet in a nn.Module
        class UNetWrapper(nn.Module):
            def __init__(self, net):
                super().__init__()
                self.net = net

            def forward(self, t, x, y=None): # Forward Pass for NeuralODE/CFM
                # t = current time step, x = input image tensor, shape (batch, channels, H, W), y = optional cond
                # Make sure input has channel dim
                if x.ndim == 3:
                    for t_j in times:
                        t_batch = t_j.expand(x1.size(0)) # vectorized over batch

                        # Sample target velocity from CFM
                        _, xt_sample, ut_sample = FM.sample_location_and_conditional_flow(xt, x1)

                        # Ensure 4D tensors
                        xt_sample = xt_sample.unsqueeze(1) if xt_sample.ndim == 3 else xt_sample
                        ut_sample = ut_sample.unsqueeze(1) if ut_sample.ndim == 3 else ut_sample

                        # Mixed precision forward
                        with autocast():
                            vt = model(t_batch, xt_sample, y)
                            step_loss = ((vt - ut_sample)**2).mean()
                            traj_loss += step_loss

                        # Euler step to update trajectory
                        xt = xt + vt * dt

                    traj_loss /= num_ode_steps
                xt = x0
```

Questions?

References

Jaiswal, S. (2025, April 5). *Multilayer Perceptrons in Machine Learning: A Comprehensive Guide*. DataCamp.

<https://www.datacamp.com/tutorial/multilayer-perceptrons-in-machine-learning> - MLP Basics

Lipman, Y., Havasi, M., Holderrieth, P., Shaul, N., Le, M., Karrer, B., Chen, R. T. Q., Lopez-Paz, D., Ben-Hamu, H., & Gat, I. (2024).

Flow Matching Guide and Code. FAIR at Meta, MIT CSAIL, Weizmann Institute of Science. <https://arxiv.org/abs/2412.06264>

Tong, A., Malkin, N., Huguet, G., Zhang, Y., Rector-Brooks, J., Fatras, K., Bengio, Y. (2023). *conditional-flow-matching: TorchCFM — Conditional Flow Matching library*. GitHub repository. <https://github.com/atong01/conditional-flow-matching> - For MLP Model

Tong, A. Y. (2023). *Flow_matching_tutorial.ipynb* [Jupyter Notebook]. GitHub. Retrieved December 1, 2025, from

https://colab.research.google.com/github/atong01/conditional-flow-matching/blob/main/examples/2D_tutorials/Flow_matching_tutorial.ipynb - Main Examples

[JamesCFM.ipynb - Colab](#)